

Studio Ousia's Quiz Bowl Question Answering System at NIPS HCQA 2017

Ikuya Yamada^{1,2} Ryuji Tamaki¹ Hiroyuki Shindo³ Yoshiyasu Takefuji²

¹Studio Ousia ²Keio University ³Nara Institute of Science and Technology

Our Background

Tokyo-based tech startup focusing on NLP and deep learning

- ▶ QA ENGINE: Deep learning-based question answering system

- ◆ Winner of two competitions

- **Human-Computer QA @ NIPS 2017**
- Human-Computer QA @ NAACL 2016



- ▶ Semantic Kernel: A state-of-the-art entity linking system

- ◆ Winner of past two competitions:

- NEEL Challenge @ WWW 2015
- W-NUT Shared Task #1 @ ACL 2015



Summary of the Task


With the assistance of his chief minister, the Duc de Sully, he lowered taxes on peasantry, promoted economic recovery, and instituted a tax on the Paulette. Victor at Ivry and Arquet, he was excluded from succession by the Treaty of Nemours, but won a great victory at Coutras.



Henry IV of France


- ▶ The system is required to guess the entity that is described in the question
- ▶ A question is given one word at a time, and the system can output an answer at any time
- ▶ The answer must be an entity that exists in Wikipedia

Overview of Our Approach



- ▶ Two *neural network models* and *information retrieval models* are combined using point-wise learning to rank model
- ▶ Our system outputs an answer if the relevance score of a top answer exceeds a threshold

Neural Quiz Bowl Solver




- ▶ *Neural Quiz Bowl Solver* solves the task as a text classification problem over possible answers with words and entities in a question as inputs
- ▶ Entities are automatically detected from a question using an entity linking method based on a dictionary built from Wikipedia
- ▶ Words, entities, and answers are mapped into continuous vectors which are initialized using Entity-Vector (Yamada et al, 2016)
- ▶ The representation of a question is computed based on the average of the word vectors (\mathbf{a}_w) and the entity vectors (\mathbf{b}_e)

$$\mathbf{v}_{D_w} = \frac{1}{N} \sum_{n=1}^N \mathbf{W}_w \mathbf{a}_{w_n}, \quad \mathbf{v}_{D_e} = \frac{1}{K} \sum_{n=1}^K \mathbf{W}_e \mathbf{b}_{e_n}$$
$$\mathbf{v}_D = \mathbf{v}_{D_w} + \mathbf{v}_{D_e}$$

- ▶ The probability of each answer being correct is computed based on the dot product between the representation of the question (\mathbf{v}_D) and the answer vector (\mathbf{c}_e):

$$P(e_t|D) = \frac{\exp(\mathbf{c}_{e_t}^\top \mathbf{v}_D)}{\sum_{e' \in \Gamma} \exp(\mathbf{c}_{e'}^\top \mathbf{v}_D)}$$

Neural Entity Type Predictor



- ▶ *Neural Entity Type Predictor* predicts the entity types of the answer (e.g., *person, author*) given a question
- ▶ Model: multi-class text classification using convolutional neural networks (CNN) (Kim, 2014)
- ▶ Target labels: entity types (automatically assigned to each answer using FIGER entity type set (Ling and Weld, 2012))

Training of Neural Network Models


- ▶ Dataset: 101,043 question-answer pairs created based on the dataset provided by competition organizers
- ▶ To perform accurately for incomplete questions, a question is truncated at a random position before inputting it to the models during training
- ▶ Stochastic gradient descent (SGD) on GPU is used to train the models

Information Retrieval Models



- ▶ Multiple relevance scores are computed against the documents associated to the answer
- ▶ Query: words in a question
- ▶ Target documents:
 - ◆ *Wikipedia text*: Page text in the answer's Wikipedia entry
 - ◆ *Questions in the dataset*: Questions associated to the answer in our dataset
- ▶ Scoring methods:
 - ◆ Okapi BM25
 - ◆ The number of common words between the question and document

Supervised Point-wise Learning to Rank



- ▶ The learning to rank model assigns a relevance score to each answer
- ▶ Algorithm: *Gradient boosted regression trees (GBRT)*
- ▶ Features: the outputs of the neural network models and the IR models
- ▶ The system outputs an answer if the relevance score of a top answer exceeds a threshold, which is set as 0.6
- ▶ For each question, we generate five questions truncated at random positions to maintain the accuracy for incomplete questions

THANK YOU!